

Tutorial 4: Value Function Approximation

Reinforcement Learning Course

Learning Objectives

By the end of this tutorial, you should be able to:

- Apply stochastic gradient descent (SGD) to value function learning
- Implement semi-gradient TD learning with neural networks
- Understand experience replay and target networks for stability
- Train a Deep Q-Network (DQN) agent on CartPole

Notation Reference

- $\hat{v}(s, \mathbf{w})$: Approximate value function with parameters \mathbf{w}
- $\hat{q}(s, a, \mathbf{w})$: Approximate action-value function with parameters \mathbf{w}
- $\mathbf{w} \in \mathbb{R}^d$: Weight vector (parameters)
- α : Learning rate (step size)
- U_t : Target value for update
- $\nabla_{\mathbf{w}}$: Gradient with respect to parameters \mathbf{w}
- \mathcal{D} : Experience replay buffer
- \mathbf{w}^- : Target network parameters

1 Stochastic Gradient Descent for RL

1.1 Deriving SGD Updates

We want to minimize the mean squared value error: $\overline{VE}(\mathbf{w}) = \mathbb{E}_{\mu}[(v_{\pi}(S) - \hat{v}(S, \mathbf{w}))^2]$

- (a) Derive the gradient descent update rule for minimizing this objective.
- (b) We don't know the true value $v_{\pi}(S)$. In Monte Carlo, we replace it with the return G_t . Write the MC SGD update.
- (c) In TD(0), we use the TD target $R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w})$ instead. Write the TD(0) SGD update.
- (d) Why is TD(0) called a "semi-gradient" method? What part of the gradient is being ignored?

2 Deep Q-Networks (DQN)

2.1 Understanding DQN Components

DQN introduced several key innovations to stabilize deep RL. For each component below, explain what problem it solves.

- (a) **Experience Replay:** Store transitions (s_t, a_t, r_t, s_{t+1}) in a replay buffer \mathcal{D} and sample random mini-batches for updates.
- (b) **Target Network:** Maintain a separate network $Q(s, a; \mathbf{w}^-)$ with frozen parameters that are updated periodically (e.g., every 1000 steps).
- (c) **ϵ -greedy Exploration:** Select random actions with probability ϵ , greedy actions otherwise.

3 Programming Exercises

In these exercises, we'll use the CartPole-v1 environment from Gymnasium. Full documentation is available at: https://gymnasium.farama.org/environments/classic_control/cart_pole/

Environment Description: A pole is attached to a cart moving along a frictionless track. The goal is to balance the pole upright by applying forces to move the cart left or right. An episode ends when:

- The pole angle exceeds ± 12 from vertical
- The cart position exceeds ± 2.4 units from center
- Episode length reaches 500 steps (solved!)

State Space: The observation is a 4-dimensional continuous vector:

Index	Variable	Range
0	Cart position (x)	$[-4.8, 4.8]$
1	Cart velocity (\dot{x})	$[-\infty, \infty]$
2	Pole angle (θ)	$[-0.418, 0.418]$ rad ($\approx \pm 24$)
3	Pole angular velocity ($\dot{\theta}$)	$[-\infty, \infty]$

Action Space: Discrete actions:

- Action 0: Push cart to the left
- Action 1: Push cart to the right

Rewards:

- +1 for every timestep the pole remains upright
- The task is considered solved when average reward ≥ 195 over 100 consecutive episodes

PyTorch Usage: This implementation uses PyTorch for building and training neural networks. If you are unfamiliar with PyTorch basics, refer to the official tutorials: <https://pytorch.org/tutorials/beginner/basics/intro.html>

3.1 Implementing Semi-Gradient TD with Linear Approximation

Open `tutorial_template.py`. First, you'll implement TD learning with linear function approximation on CartPole with hand-crafted features.

CartPole state: $[x, \dot{x}, \theta, \dot{\theta}]$ where x is cart position, θ is pole angle.

- (a) **Implement `create_features`:** Create a feature vector from the state. Use polynomial features up to degree 2 (including cross terms).
- (b) **Implement `LinearQNetwork`:** A linear Q-network where $Q(s, a) = \mathbf{w}_a^\top \phi(s)$ with separate weights for each action.
- (c) **Implement `update`:** The semi-gradient TD update for linear Q-learning:

$$\mathbf{w}_a \leftarrow \mathbf{w}_a + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)] \phi(s)$$

- (d) **Test your implementation:** Run the training loop. The agent should solve CartPole (average reward > 195 over 100 episodes) within 500-1000 episodes.

Expected output:

- Training should show gradual improvement
- Episodes will be short initially (pole falls quickly)
- After convergence, episodes should reach maximum length (500 steps)
- Plot shows increasing average reward over episodes

3.2 Implementing Deep Q-Network (DQN)

Now implement the full DQN algorithm with experience replay and target networks.

- (a) **Implement `ReplayBuffer`:** Store transitions and sample random mini-batches.
 - Store tuples $(s_t, a_t, r_t, s_{t+1}, \text{done}_t)$
 - Implement `add()` and `sample(batch_size)`
- (b) **Implement `QNetwork`:** A neural network with:
 - Input: state vector (4 dimensions for CartPole)
 - Hidden layers: 2 layers with 128 units each, ReLU activation
 - Output: Q-values for each action (2 for CartPole)
- (c) **Implement the DQN training loop:**
 - Select actions using ϵ -greedy policy
 - Store transitions in replay buffer
 - Sample mini-batches and compute TD targets using target network
 - Update Q-network using MSE loss
 - Periodically update target network: $\mathbf{w}^- \leftarrow \mathbf{w}$

(d) **Hyperparameters to use:**

- Batch size: 64
- Learning rate: 0.001
- γ : 0.99
- ϵ : start at 1.0, decay to 0.01
- Target network update frequency: every 100 steps
- Replay buffer size: 10,000

Expected output:

- Your code should generate plots showing training progress
- DQN should solve CartPole faster than linear approximation (200-400 episodes)
- Learning curve should be smoother than linear due to replay buffer
- Final policy should consistently balance the pole for 500 steps

4 Reflection Questions

4.1 Comparing Methods

After implementing both linear and deep Q-learning:

- Which method learned faster? Why?
- Which method achieved better final performance? Why?
- What are the advantages and disadvantages of each approach?
- When would you choose linear function approximation over deep learning?